

ECE 671

Introduction to

Computer Networks

Week 3 Lesson 5

Transport Layer & Congestion Control

Rationale

- **Transport layer is an important part of Internet protocol stack**
- **Understanding complete protocol stack is important**
- **Transport layer services are used by application layer**
 - **Even if someone does not care about networking details, understanding which transport layer protocol to use (and why) is really important**
- **Congestion control in transport layer shows how resource sharing is achieved without central control**

Objectives

- Explain the process of connection setup and teardown in TCP
- Evaluate causes of congestion and possible solutions in a transport layer
- Analyze TCP approaches to mitigate congestion

Prior Knowledge


- **Transport layer connect application processes**
 - **Different transport layer protocols depending on requirements**
- **Transport layer features in last lesson**
 - **De-/multiplexing**
 - **Reliability**

Orchestrated Discussion (Hand Raise): Lesson Reflection Feedback

- Discuss questions and comments on Lesson Reflection from prior lesson


TCP Connection Management

- Connection setup
 - Three-way handshake
 - SYN
 - SYN/ACK
 - ACK
 - SYN counts as one byte
 - ACK may carry data already
 - Flag in header identifies SYN
 - Used in network systems to identify new connections



TCP Connection Management

- Connection teardown
 - Each side closes when transmission is complete
 - FIN
 - ACK
 - Final FIN or ACK may get lost
 - Need to be able to retransmit
- Connection cleanup after timed wait



Document Cam: Performance Expectations

- Read “For Impatient Web Users, an Eye Blink Is Just Too Long to Wait” (by Steve Lohr, New York Times, 2/29/2012)

Group Discussion and Report Back (Short Answer): Performance Expectations

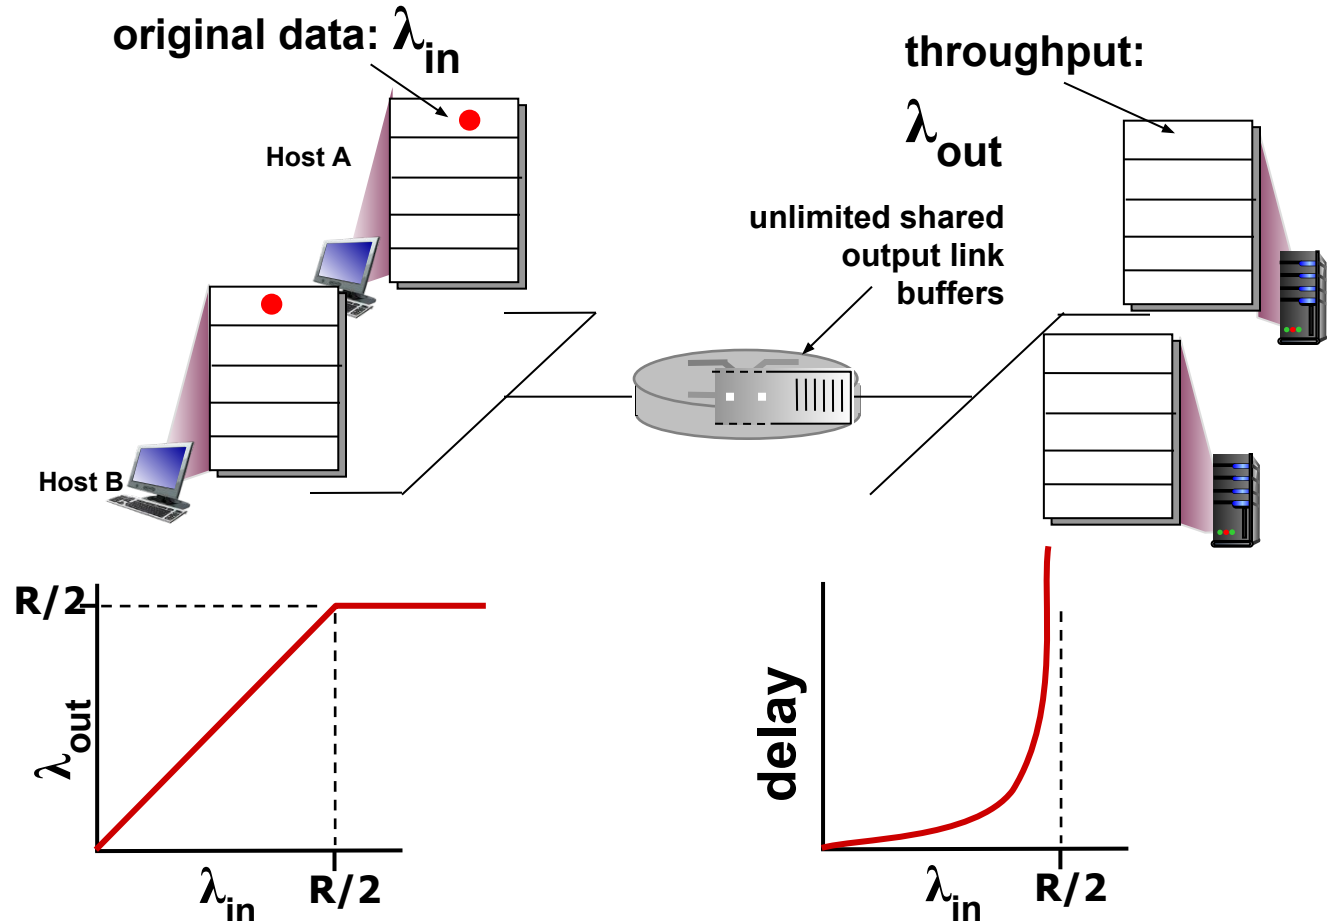
- Answer the following questions:
 - What are users' expectations for Internet interactions?
 - What are possible ways for providers to satisfy users?

Network Congestion

- What is the problem when all providers send always at “full speed”?
 - Network resources are limited
 - Network links will be fully utilized
 - Router buffers may fill up and drop packets
- Network congestion occurs when too much traffic is sent by end-systems
 - Can be localized: congested link
 - Can be larger: multiple congested links and routers

Whiteboard: Cause of Congestion - Scenario 1


- Baseline scenario:
 - Two senders, two receivers
 - one router, infinite buffers
 - output link capacity: R
 - No retransmission



- maximum per-connection throughput: $R/2$
- large delays as arrival rate, λ_{in} , approaches capacity



Cause of Congestion: Scenario 2

- One router, finite buffers
- Sender retransmission of timed-out packet
 - Application-layer input = application-layer output: $\lambda_{in} = \lambda_{out}$
 - Transport-layer input includes retransmissions : $\lambda_{in}' \geq \lambda_{in}$




Cause of Congestion: Scenario 2

- Idealization: perfect knowledge
 - Sender sends only when router buffers available





Cause of Congestion: Scenario 2

- Idealization: known loss
 - Packets can be lost, dropped at router due to full buffers
 - Sender only resends if packet known to be lost



Cause of Congestion: Scenario 2



- Idealization: known loss
 - Packets can be lost, dropped at router due to full buffers
 - Sender only resends if packet known to be lost



Cause of Congestion: Scenario 2

- Realistic: duplicates

- Packets can be lost, dropped at router due to full buffers
- Sender times out prematurely, sending two copies, both of which are delivered




Cause of Congestion: Scenario 2

- **Realistic: duplicates**

- Packets can be lost, dropped at router due to full buffers
- Sender times out prematurely, sending two copies, both of which are delivered

- “Costs” of congestion:

- More work (retransmission) for given “goodput”
- Unneeded retransmissions: link carries multiple copies of packets
- Decreasing goodput




Cause of Congestion: Scenario 3

Q: What happens as λ_{in} and λ_{in}' increase?



• Scenario:

- Four senders
- Multihop paths
- Timeout/retransmit

A: As red λ_{in}' increases, all arriving blue packets at upper queue are dropped, blue throughput $\rightarrow 0$



Cause of Congestion: Scenario 3



- Another “cost” of congestion:
 - When packet dropped, any “upstream transmission capacity used for that packet was wasted!

Orchestrated Discussion (Hand Raise): Solving the Congestion Problem


- What can we do about the congestion problem?
 - We will brainstorm as a class

Congestion Control Approaches

- Reserving resources (cf. circuit switching)
 - As for permission before sending
 - Send at permitted rate
 - Congestion can be avoided through managing reservations
- Adapting sending rate (cf. packet switching)
 - Detect congestion (e.g., packet losses cause NAKs)
 - Reduce sending rate to avoid congestion collapse
- Tradeoff: central control vs. distributed approach


TCP Congestion Control

- Additive increase, multiplicative decrease (AIMD)
- Approach: sender increases transmission rate (window size), probing for usable bandwidth, until loss occurs
 - Additive increase: increase `cwnd` by 1 MSS every RTT until loss detected
 - Multiplicative decrease: cut `cwnd` in half after loss



AIMD saw tooth
behavior: probing
for bandwidth

TCP Congestion Control



- **Sender limits transmission:**

$$\text{LastByteSent} - \text{LastByteAcked} \leq \text{cwnd}$$
- **cwnd** is dynamic, function of perceived network congestion


TCP sending rate:

- Roughly: send cwnd bytes, wait RTT for ACKS, then send more bytes

$$\text{rate} \approx \frac{\text{cwnd}}{\text{RTT}} \text{ bytes/sec}$$

TCP Slow Start

- When connection begins, increase rate exponentially until first loss event:
 - Initially `cwnd` = 1 MSS
 - Double `cwnd` every RTT
 - Done by incrementing `cwnd` for every ACK received
- Summary: initial rate is slow but ramps up exponentially fast



TCP Loss Detection and Reaction

- Loss indicated by timeout:
 - **cwnd** set to 1 MSS;
 - Window then grows exponentially (as in slow start) to threshold, then grows linearly
- Loss indicated by 3 duplicate ACKs: TCP RENO
 - Duplicate ACKs indicate the network is capable of delivering some segments
 - **cwnd** is cut in half, window then grows linearly
- TCP Tahoe always sets **cwnd** to 1 (timeout or 3 duplicate ACKs)


TCP Slow Start and Congestion Avoidance

Q: When should the exponential increase switch to linear?

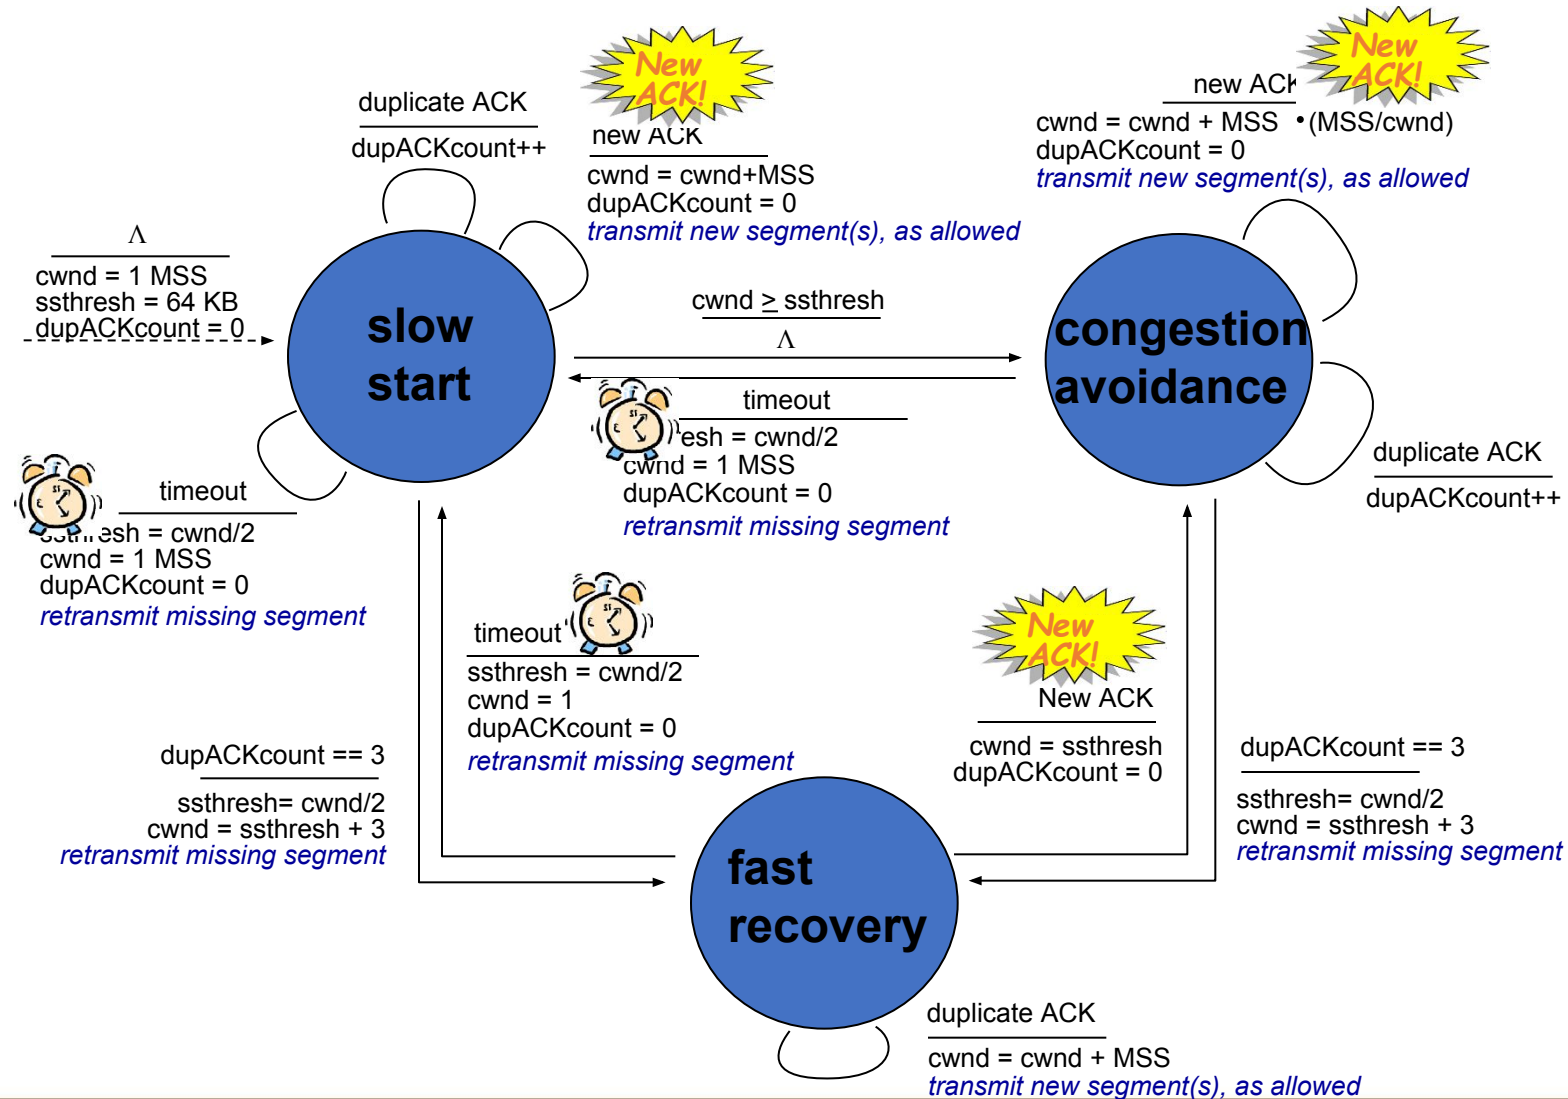
A: When cwnd gets to 1/2 of its value before timeout.

Implementation:

- Variable ssthresh
- On loss event, ssthresh is set to 1/2 of cwnd just before loss event




TCP Congestion Control Summary




Poll: TCP Throughput

- What is the average TCP throughput as a function of window size W and RTT?
- Ignore slow start, assume always data to send



TCP Throughput


- What is the average TCP throughput as function of window size W and RTT?
 - Ignore slow start, assume always data to send
- W : window size (measured in bytes) where loss occurs
 - Average window size (# in-flight bytes) is $\frac{3}{4} W$
 - Average throughput is $\frac{3}{4}W$ per RTT



$$\text{Avg. TCP throughput} = \frac{3}{4} \frac{W}{\text{RTT}} \text{ bytes/sec}$$

Link Sharing with TCP


- Two TCP connections share a link:



- Eventually each connection receives a fair share
 - How can this be shown?


Link Sharing with TCP

- Illustration of two connections



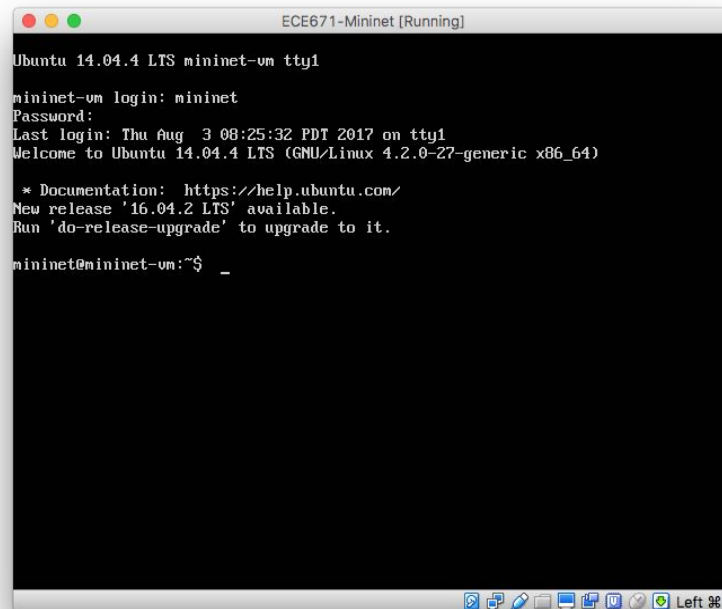
Link Sharing with TCP

- Illustration of two connections



Connected Device: HTTP & TCP Demo

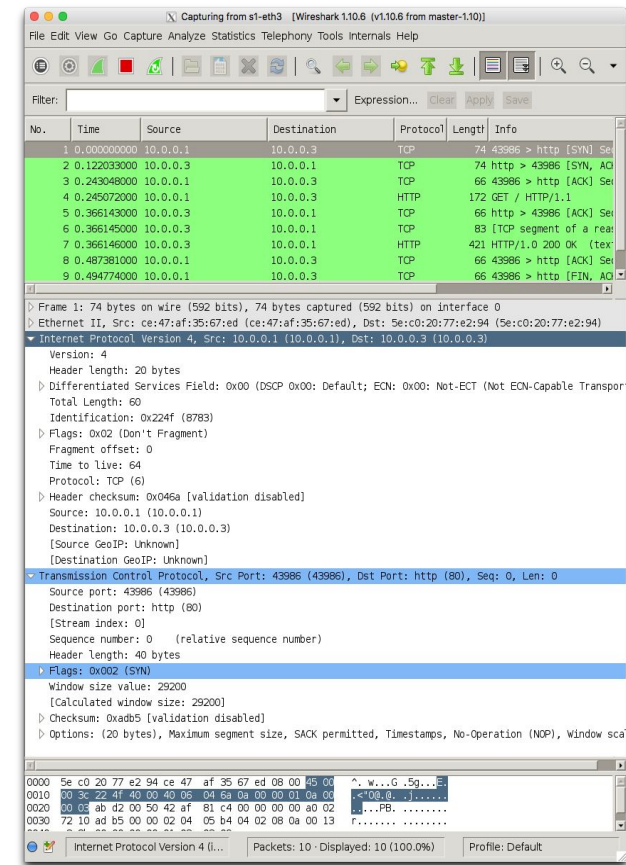
- Mininet demonstration in virtual machine



```
ECE671-Mininet [Running]
Ubuntu 14.04.4 LTS mininet-vm tty1
mininet-vm login: mininet
Password:
Last login: Thu Aug 3 08:25:32 PDT 2017 on tty1
Welcome to Ubuntu 14.04.4 LTS (GNU/Linux 4.2.0-27-generic x86_64)

 * Documentation:  https://help.ubuntu.com/
New release '16.04.2 LTS' available.
Run 'do-release-upgrade' to upgrade to it.

mininet@mininet-vm:~$ _
```



Summary of Lesson

- Connection setup and teardown
- Reasons for congestion
- Congestion control in TCP
- Fairness between two connections
- Demo

Post-work for Lesson 5

Homework #3

- After the Live Lecture, you will complete and submit a homework assignment. Go to the online classroom to view and submit the assignment.

To Prepare for the Next Lesson

- Complete and submit the Post-work for Lesson 5.
- Read the Required Readings for Lesson 6.
- Complete the Pre-work for Lesson 6.

Go to the online classroom for details.